

# Executive Brief — Shunyaya Symbolic Mathematical AI (SSM-AI)

A conservative, drop-in alignment lane that shows stability and reduces surprises — **without changing your numbers.**

Status: Public Research Release (v1.8)

Date: 27 October 2025

Caution: Research/observation only. Not for critical decision-making.

---

## Three business questions your leadership should be able to answer

1. **Can you see the cliff before revenue falls?** When KPIs look great, does a stability lane flag growing tail risk early?
  2. **Are weekly wins actually durable?** When totals tie, can you see drift that predicts next month's miss or overage?
  3. **Are you paying a stability tax?** When a unit looks efficient, does a band view expose retries, latency tails, and incident risk you are not pricing?
- 

## Three AI platform questions your product teams should be able to answer

1. **Are model choices sturdy under real traffic?** When two vendors score similarly on accuracy, does a bounded chooser  $RSI \text{ in } (-1, +1)$  reveal which one stays stable through spikes, tools, and RAG drift?
  2. **Do you catch failure faster than you ship it?** Can you band and gate alignment ( $RSI_{env} := g_t * RSI$ ) to quarantine regressions before they hit users or costs?
  3. **Can you replay decisions exactly?** If a ranking or route changes, can you reconstruct it from stamped  $U/W$  roll-ups — order-invariant and vendor-neutral — without touching your original metrics ( $\phi((m, a)) = m$ )?
- 

## What it is (one line)

A conservative add-on that leaves every KPI exactly as is and adds a small, bounded stability lane beside it:  $x := (m, a)$  with  $a \text{ in } (-1, +1)$  and collapse parity  $\phi((m, a)) = m$ ; decisions read a simple chooser  $RSI \text{ in } (-1, +1)$  and bands  $(A++/A+/A0/A-/A--)$  to show sturdiness — no rewrites, no PII, and replayable stamps if you want them.

---

## Why leaders care today (outcomes over math)

- **Fewer surprises, earlier.** Stability bands surface tail risks and drift before classical KPIs move.
  - **Cleaner vendor choices.** A single bounded chooser  $RSI$  in  $(-1, +1)$  makes cross-model comparisons fair; no calibration wars.
  - **Lower run costs without guesswork.** Token, latency, and retry improvements show up as stable  $A+/A++$  bands; savings remain outside the lane math.
  - **Quieter operations.** Band hysteresis trims alert flicker; teams act on trends, not noise.
  - **Evidence you can replay.** Optional stamps and `knobs_hash` make decisions tamper-evident and reproducible, end to end.
- 

## What SSM-AI delivers (at a glance)

- **Confidence in every output.** Whether it is a model response, a forecast, or a KPI, each carries a visible stability lane that signals reliability in real time.
  - **Vendor-neutral truth layer.** Adds bounded symbolic entropy  $(RSI, a)$  over existing outputs, so teams can compare models and tools without bias or retraining.
  - **Zero friction deployment.** Drop-in SDK; no model retrain, no data migration, and no infrastructure rewrite.
  - **Faster AI economics.** Drift detection and early failure cues help cut wasted tokens, retries, and review loops — all while leaving the original metrics intact  $(\phi((m, a)) = m)$ .
  - **Ethical and auditable by design.** Optional SSM-Clock stamps and `knobs_hash` provide full replayability, offering audit confidence across the AI value chain.
- 

## Under the hood (roots, one-liners)

- **Core math (Shunyaya Symbolic Mathematics — SSM).** Two-lane numeral  $x := (m, a)$  with  $a$  in  $(-1, +1)$ . Safe map:  $\text{clamp} \rightarrow u := \text{atanh}(a) \rightarrow \text{sum in } u \rightarrow a := \tanh(u)$ . Collapse parity holds:  $\phi((m, a)) = m$ .
- **Symbols (SSMS).** Small verb set for clean connectors: `CLAMP`, `MAP(atanh/tanh)`, `FUSE(U/W)`, `MUL_DIV(M2)`, `BAND`, `GATE`, `STAMP`.
- **Time & provenance (SSM-Clock + SSM-Clock Stamp).** Optional, tamper-evident stamps and order-invariant roll-ups in  $u$ -space. One-line stamp:  
`SSMCLOCK1|iso_utc|rasi_idx|theta_deg|sha256(file)|chain.`
- **Hardware (SSMH).** Same lane math on tiny fixed-point MACs: `atanh-in`, `tanh-out`, streaming  $(U, W)$  accumulators for deterministic, low-cost speedups.
- **Search (SSM-Search).** Lane-native retrieval: rank by bounded  $RSI$  while classical retrieval scores remain untouched.

---

## What is included in SSM-AI (integration now).

Included today: SSM core lane math, SSMS verbs, SSM-Search scorer/ranker, SSM-Audit KPI overlays, optional SSM-Clock Stamp fields, and the SSM-Hardware parity spec (software↔fixed-point equivalence).

### Roadmap for SSM-AI integration.

Domain-specific adapters delivered as manifest presets and evidence packs (observation-only), built on SSM modules — for example, an SSM-Chemistry adapter.

**One-liner.** One bounded lane, many surfaces: `clamp` → `atanh` → `sum` → `tanh` powers provenance, acceleration, search, and audit overlays — while `phi((m,a)) = m` keeps the original numbers pristine.

---

## Limits & failure modes (candor box)

- **Low-signal regimes.** If the declared lens is weak or noisy, `a` will hover near 0; treat bands as “insufficient evidence,” not a verdict.
  - **Poor lenses.** Badly designed lenses mislead any chooser. Publish lens recipes and stick to bounded transforms.
  - **Throughput before acceleration.** Extremely high-throughput paths should use SSMH or vectorized u-fuse to avoid CPU hotspots.
  - **Division policy.** Defaults are strict; if you need meadow-style semantics, declare it explicitly and test with acceptance checks.
- 

## In one line

SSM-AI makes stability visible and measurable — so your AI runs calmer, cheaper, and more trustworthy, with zero code forks.

---

## Engineer’s 1-Pager — Evaluation Guide

**Kernel (never touches m).**

- **Object:** `x := (m, a)` with `a` in `(-1,+1)`; **collapse:** `phi((m,a)) = m`.
- **Clamp:** `a_c := clamp(a, -1+eps_a, +1-eps_a)`.
- **Map:** `u := atanh(a_c)`; **inverse** `a := tanh(u)`.
- **Fuse (order-invariant):** `U += w*atanh(a)`; `W += w`; `a_out := tanh( U / max(W, eps_w) )`.
- **Weights:** `w := |m|^gamma` (default `gamma = 1`).

- **Mul/Div (M2 by default):**

- **Multiply:**  $(m1, a1) * (m2, a2) := (m1 * m2, \tanh(\operatorname{atanh}(a1\_c) + \operatorname{atanh}(a2\_c)))$
- **Divide (strict,  $m2 \neq 0$ ):**  $(m1, a1) / (m2, a2) := (m1 / m2, \tanh(\operatorname{atanh}(a1\_c) - \operatorname{atanh}(a2\_c)))$

## Chooser & gating.

- **Chooser:**  $RSI := \tanh((V_{out} - U_{in}) / \max(W_{in}, \epsilon_w))$  (declare lens  $\rightarrow$  produce  $U_{in}, V_{out}, W_{in}$ ).
- **Calm gate (optional):**  $RSI_{env} := g_t * RSI$ , with  $g_t$  in  $[0, 1]$  built from bounded drift lanes; keep gate observation-only.

## Defaults (normative, copy into manifest).

- $\epsilon_a = 1e-6, \epsilon_w = 1e-12, \gamma = 1$ .
- **Bands:** A++/A+/A0/A-/A-- (publish your numeric thresholds once).
- **Division policy:** strict by default; alternatives must be declared and tested.
- **Rounding:** half-even at export,  $\geq 6$  decimals; full precision internal.

## 12-line pseudocode (reference).

```
a_c = clamp(a, -1+eps_a, +1-eps_a)
u  = atanh(a_c)
for each stream item:
    w = abs(m)**gamma
    U += w * atanh(clamp(a_i, -1+eps_a, +1-eps_a))
    W += w
a_out = tanh( U / max(W, eps_w) )
RSI    = tanh( (V_out - U_in) / max(W_in, eps_w) )
RSI_env = g_t * RSI    # optional
band   = band_of(a_out) # A+..A--
assert phi((m, a_out)) == m
```

## Manifest stub (minimal).

```
{
  "eps_a": 1e-6, "eps_w": 1e-12, "gamma": 1,
  "bands": { "A++": 0.80, "A+": 0.40, "A0": -0.10, "A-": -0.40, "A--": -
0.80 },
  "division_policy": "strict",
  "rounding": { "mode": "bankers", "decimals": 6, "when": "export" },
  "gate": { "enabled": false } # set true only if you publish Z/A recipes
```

}

---

### Acceptance checks (drop-in).

- **Collapse parity:** `phi((m, a)) == m` after every op.
- **Order-invariance:** shuffle test on any stream yields `|a_out_diff| <= 1e-12`.
- **Division policy:** strict test near `m2 -> 0` fails fast; meadow only if declared.
- **Bounds:** always enforce `|a| < 1, |RSI| < 1`.
- **Determinism:** same inputs  $\rightarrow$  identical `U`, `W`, `a_out`, `band` after rounding.

---

### Golden vectors to wire into CI (starter set).

1. Clamp edges: `a = ±1`  $\rightarrow$  clamped, no NaNs.
2. Long-path fuse: 100-step stream equals batched fuse.
3. Mul/Div (M2): `a' * = tanh(atanh(a1) +/- atanh(a2))`.
4. Bands + hysteresis: promote/demote only past declared deltas.
5. Strict division near zero: enforce policy choice.
6. Decode rerank micro (text): bounded `RSI` selects the calmer option with identical `m`.

---

### What changes in your code.

Nothing on `m`. You attach `(a, band)` beside your existing outputs; your selection logic can read `RSI` or `RSI_env` as a bounded, comparable index.

---

OMP